

A Comparison of the Use of Virtual Versus Physical Snapshots for Supporting Update-Intensive Workloads

Darius Šidlauskas
Aalborg University
darius@cs.aau.dk

Christian S. Jensen
Aarhus University
csj@cs.au.dk

Simonas Šaltenis
Aalborg University
simas@cs.aau.dk

DaMoN 2012
Scottsdale, Arizona, USA
May 21, 2012

Center for Data-intensive Systems

Motivation (1)



- Parallelism in chip multi-processors continues to increase
- Writing parallel code is difficult
 - Complications of concurrency control
 - Phantom problems
 - Deadlock situations
 - Serialization bottlenecks
 - Especially, for workloads that intermix frequent updates with queries

Motivation (2)



- A snapshot-based approach
 - Write-only snapshot (S_1) for write operations
 - Read-only snapshot (S_2) for read operations
 - Regularly, S_2 is refreshed based on S_1
- Advantages:
 - Isolates otherwise conflicting operations
 - Enables atomic (full) data scans
 - Simplifies concurrency control (and its verification)
 - Current technologies permit very frequent snapshotting

How to create a snapshot?



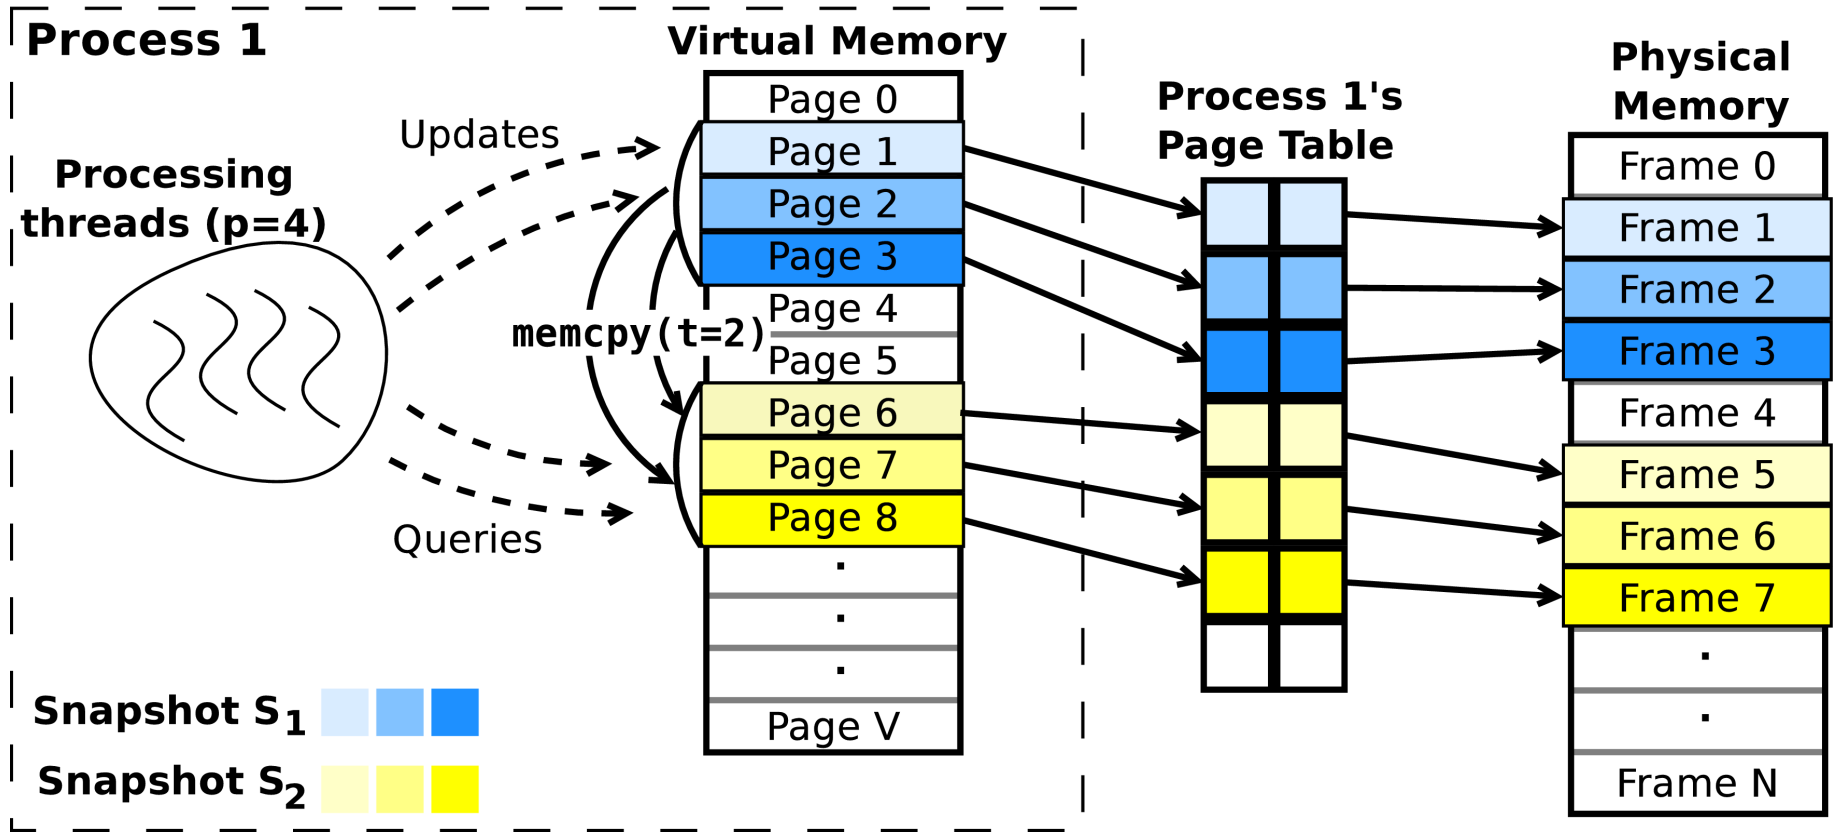
- Physical snapshots
 - `memcpy`-based
 - Full snapshots (eager copying)
 - Built-in hardware support: prefetching
 - TwinGrid (*SSTD '11*)
- Virtual Snapshots
 - `fork`-based
 - Incremental snapshots (copy-on-write)
 - Built-in hardware support: MMU, TLB
 - HyPer (*ICDE '11*)

Setting

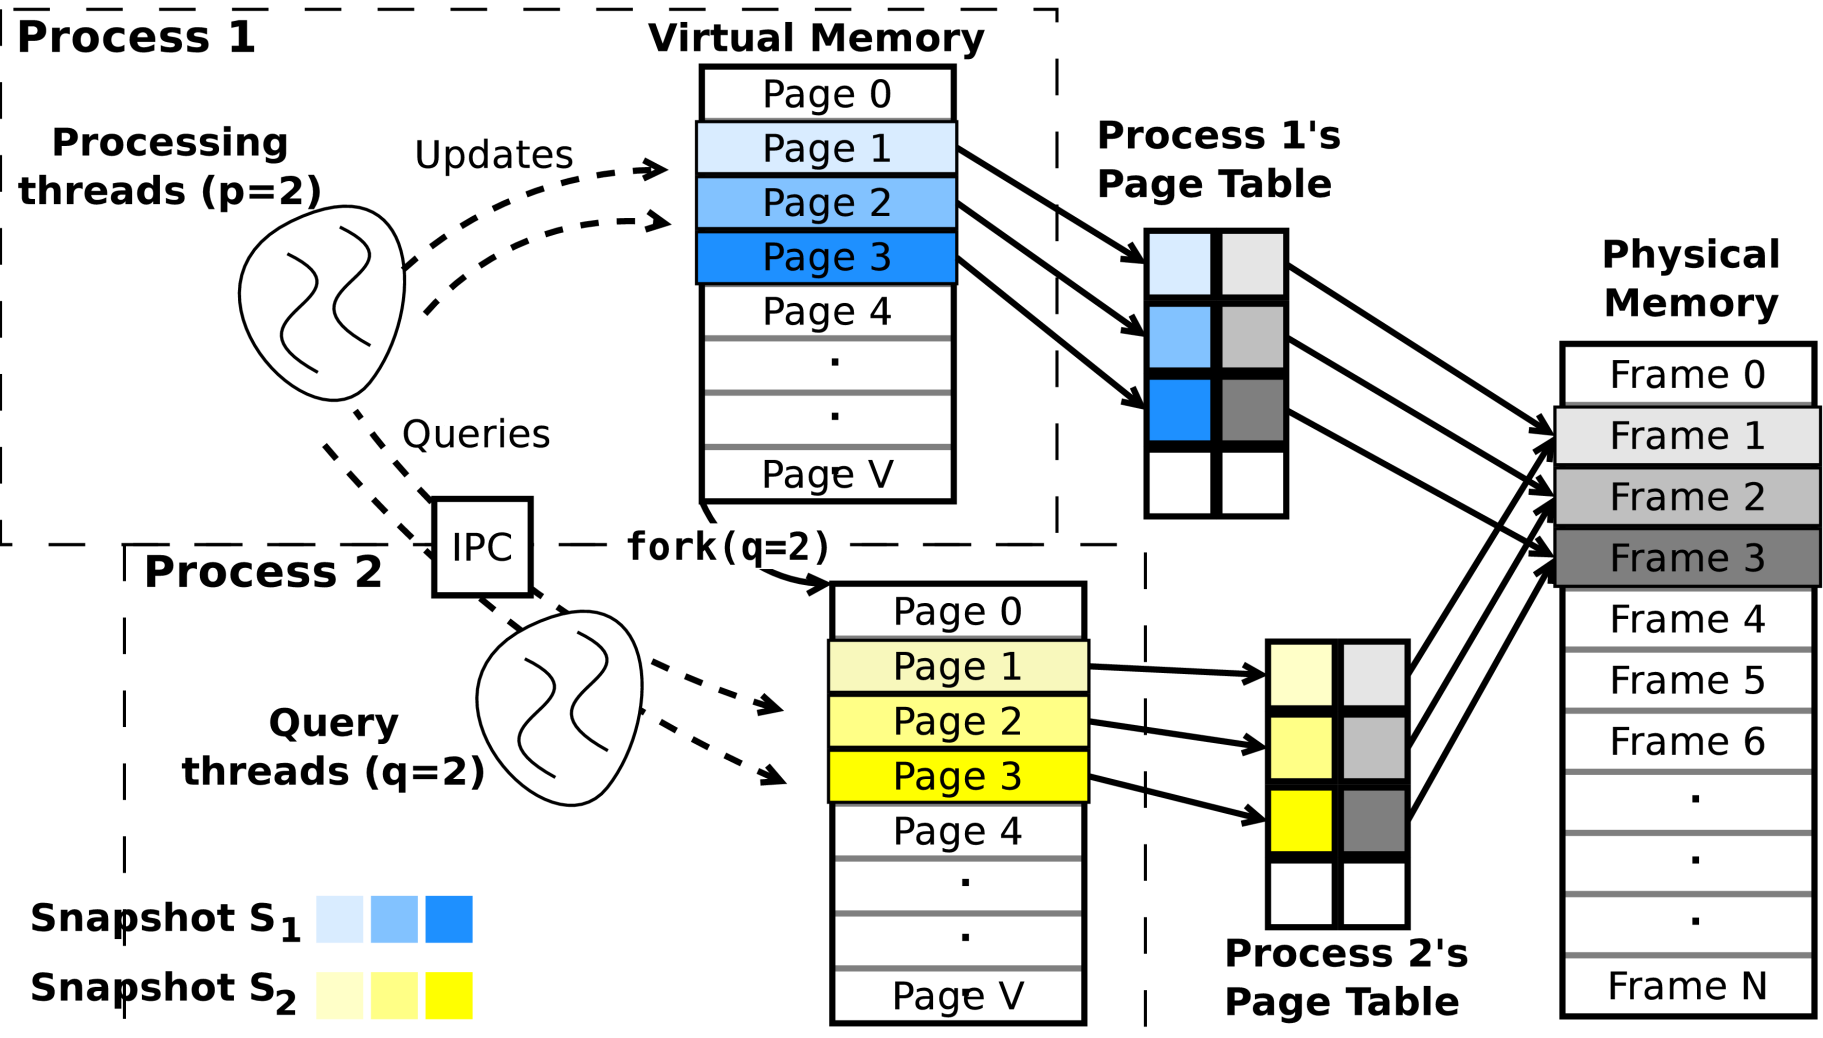


- Two snapshots
 - Each is an array of N items of fixed size (64 bits)
 - S_1 is for updates (atomic 64-bit writes)
 - S_2 is for queries (range scans)
- Snapshotting frequency (F_s)
 - Expressed as a number of updates
 - E.g., snapshotting occurs every 1000 updates
 - Controls the freshness of S_2

Physical Snapshotting



Virtual Snapshotting



Performance Study

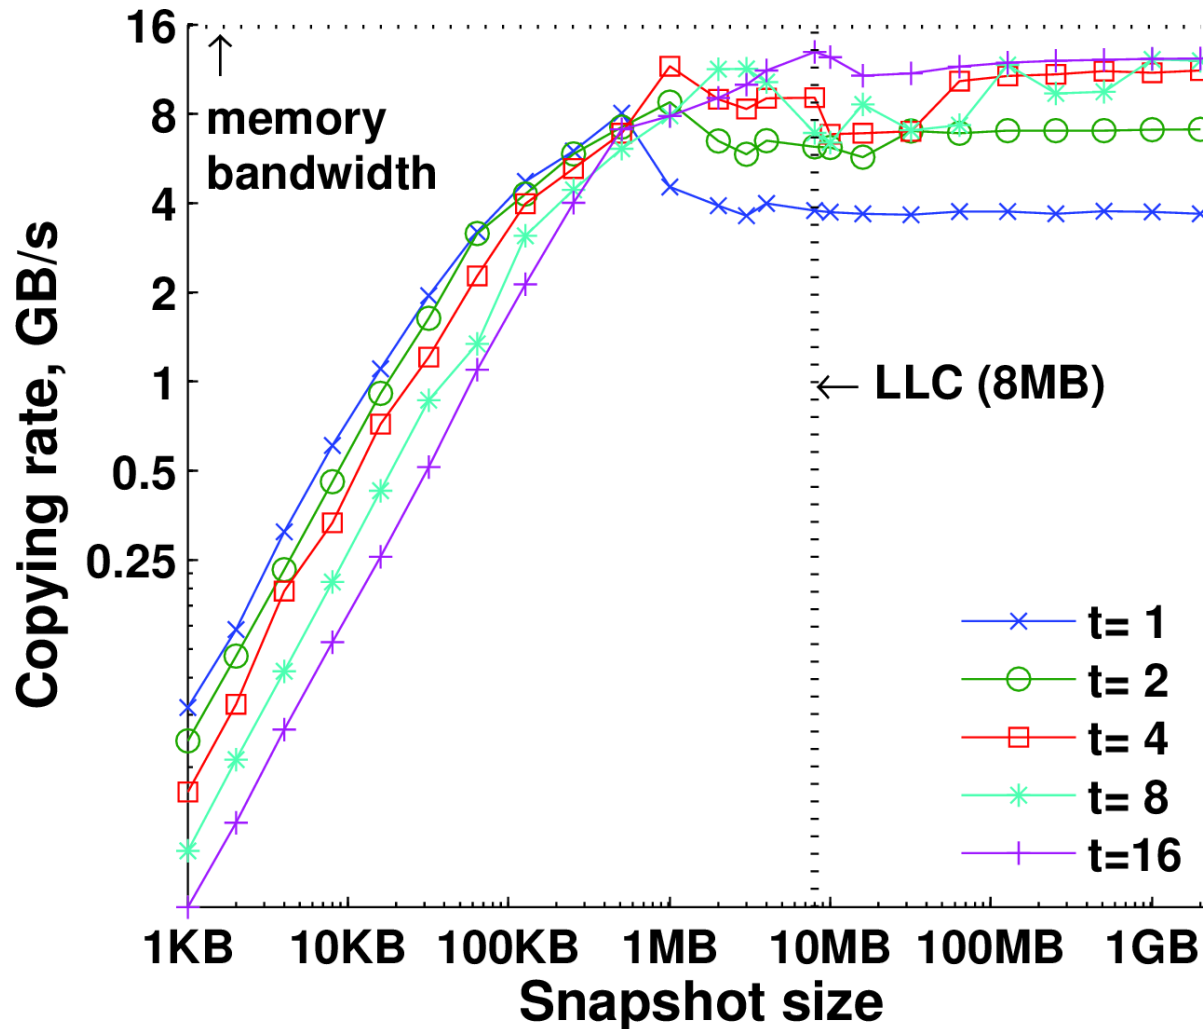


- Considers
 - Time per snapshot creation
 - Snapshot updating under different update distributions
 - Cost to query a snapshot
 - Memory consumption
- Includes four diverse multi-core platforms
 - Lean-camp and fat-camp CMP designs
 - Single- and multi-socket configurations
 - Two different Unix flavors

memcpy Performance



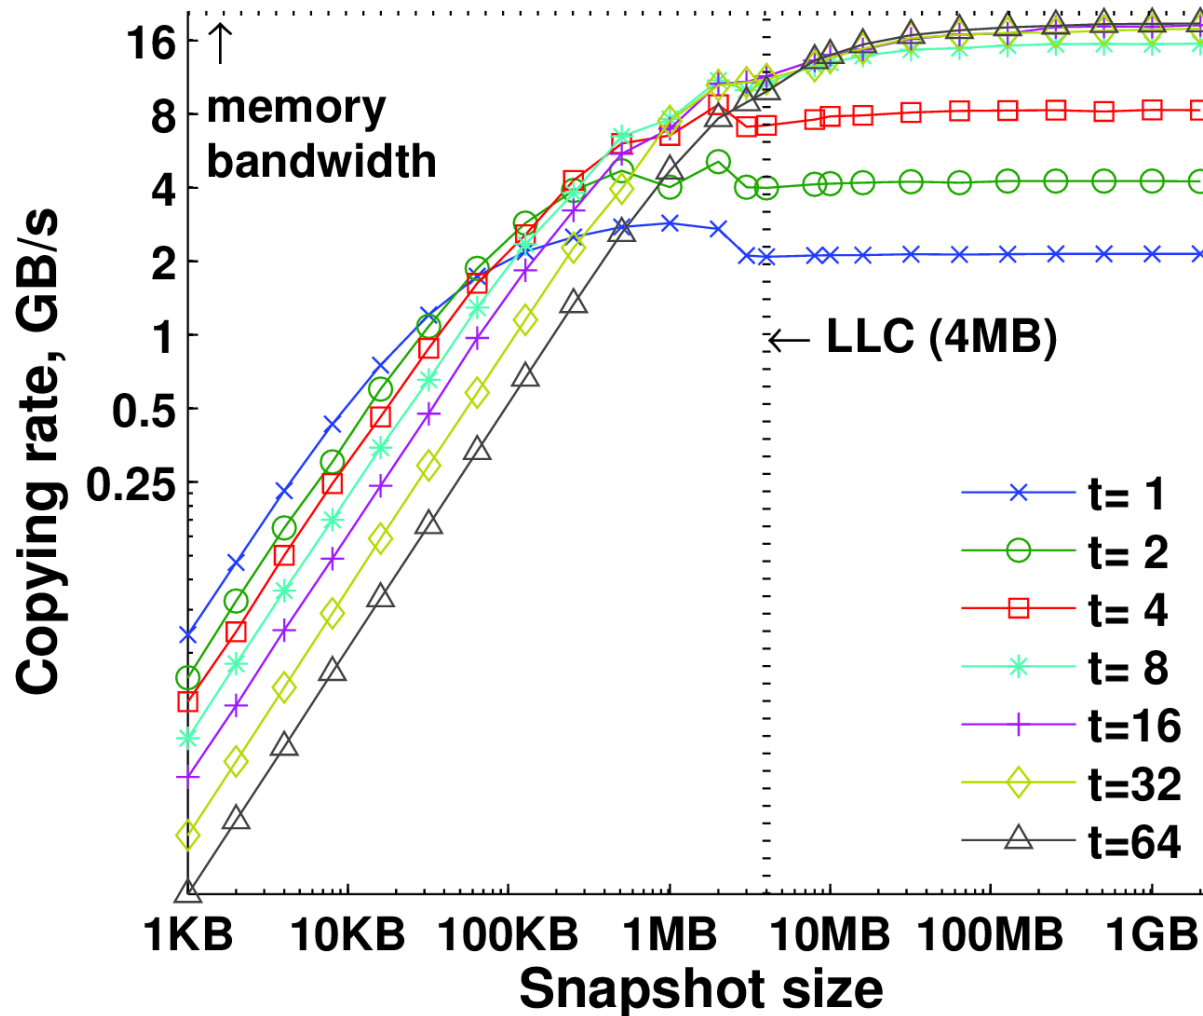
- Dual 4-core Intel Xeon X5550 2.67GHz with 16 thread contexts




memcpy Performance



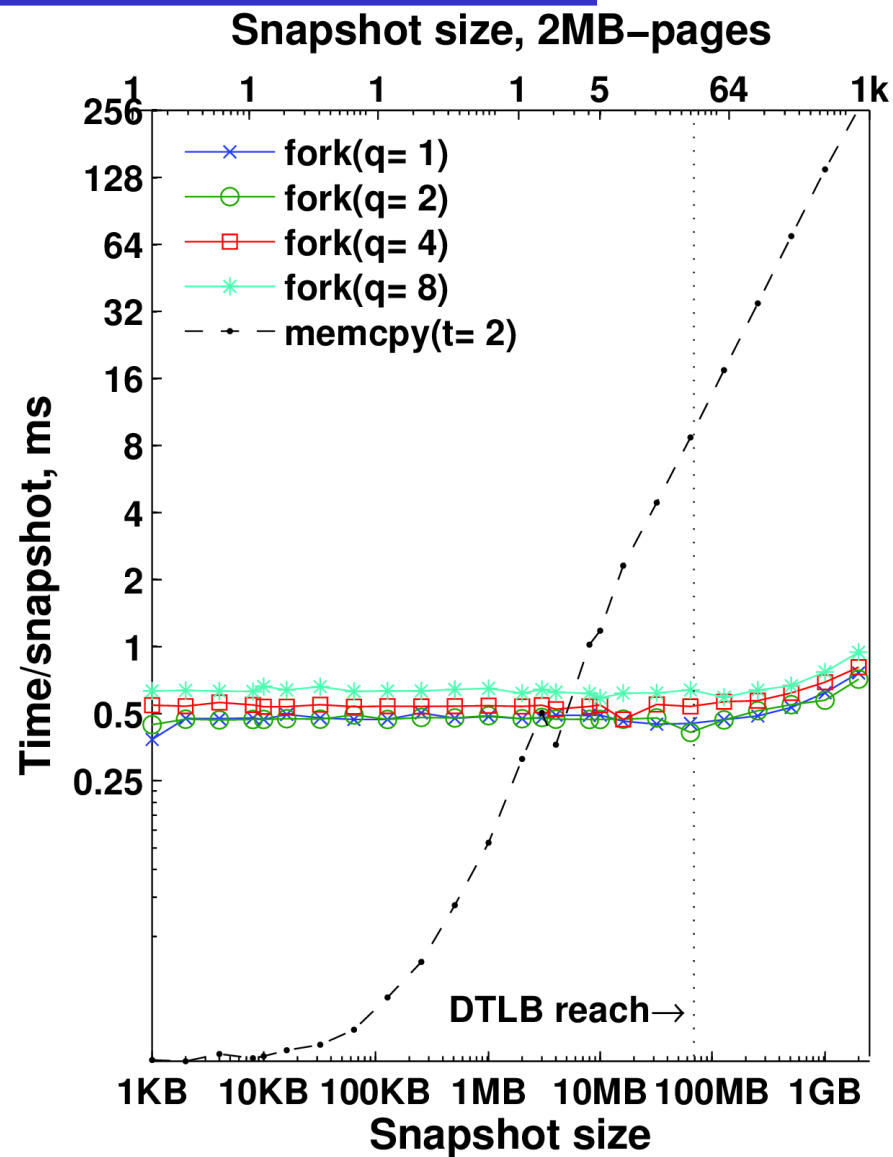
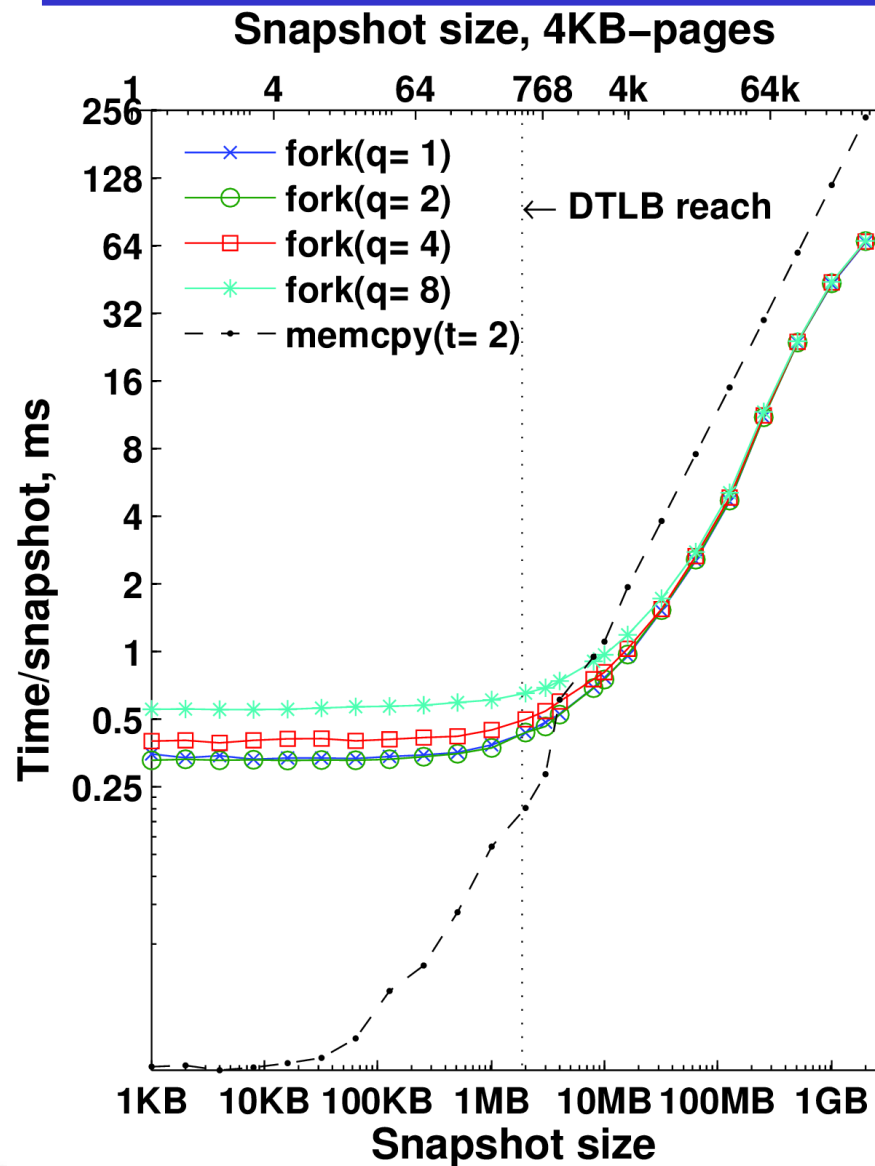
- 8-core UltraSPARC-T2 1.2GHz (Niagara 2) with 64 threads



fork VS memcpy




fork VS memcpy



Update Performance



- Snapshot size: 2^{24} items or 128MB
- Y-axis: msec per iteration (snapshotting + update processing)
- X-axis: snapshotting frequency



Conclusions



- Virtual Snapshotting
 - Implicitly supports linked data structures (pointers and non-continuous snapshots)
 - Benefits greatly from huge pages
 - Smaller memory footprint
- Physical Snapshotting
 - Easier to implement query support
 - Creation of large snapshots is very competitive
 - Efficient on all considered platform
 - Surprisingly efficient under many update distributions