

GiST Scan Acceleration Using Coprocessors

Felix Beier, Torsten Kilius, Kai-Uwe Sattler

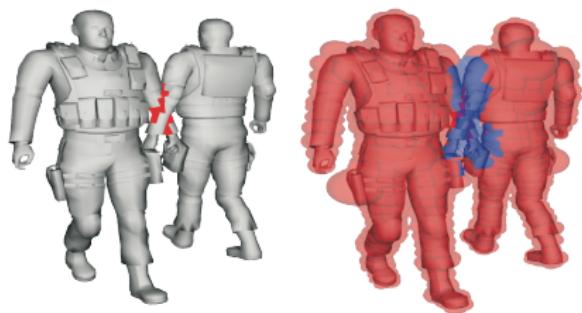
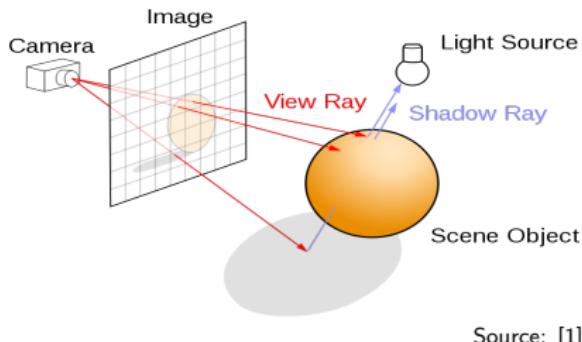
Ilmenau University of Technology

05/21/2012

Outline

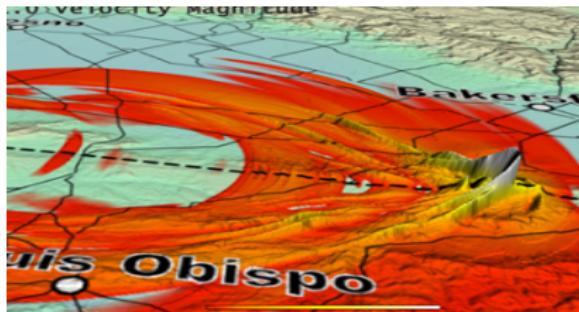
- 1 Introduction
- 2 GiST Hardware Abstraction Layer
- 3 Evaluation
- 4 Summary

Co-processing Index Searches

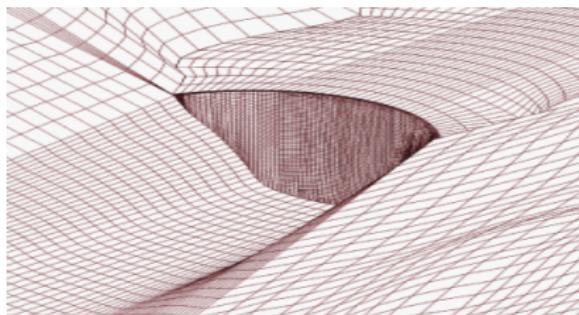


- Ray tracing: many independent point queries
 - Collision detection (spatial join): many independent range queries
 - Utilization of massive parallelism offered by modern coprocessors
- Special index structures carefully tuned for specific hardware

Index Frameworks



Source: [3]

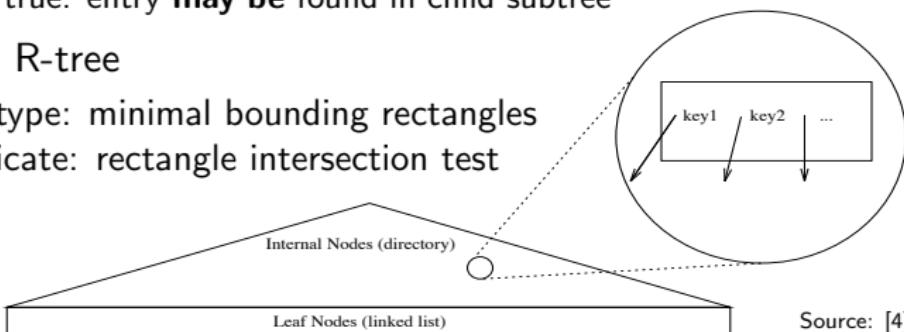


Source: [2]

- Various applications require specialized index structures
 - Scientific data
 - Enormous data volumes
 - Unknown data characteristics
 - Costly prototyping
 - Gap: scientists vs. system developers
- Rapid index development with frameworks like GiST

GiST - Generalized Index Search Tree

- Framework for implementation of height-balanced search trees
 - Implements common tree operations (insertions, deletions, node splits, height-balancing)
 - Developer specifies key data type and type-specific operations
 - Lookup predicate returns
 - false: entry can **not** be found in child subtree
 - true: entry **may be** found in child subtree
- Example: R-tree
 - Key type: minimal bounding rectangles
 - Predicate: rectangle intersection test



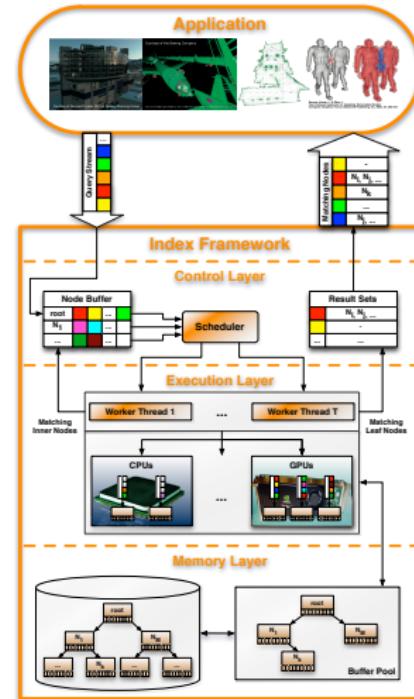
Source: [4]

Motivation

- Combining the best of both worlds
 - Extensibility of GiST framework
 - Performance improvements through co-processing
- Challenges
 - Finding fine-grained parallel algorithms
 - Utilizing hardware capabilities
 - Out-of-core implementation
 - Consideration of co-processing overheads

Framework Design

- Applications issue stream of query batches
- Iterator for matching leaf nodes
- Grouping queries to node batches for better locality
- Specialized scan implementation for various (co)processors
- Automatic scheduling to best execution unit
- Out-of-core implementation



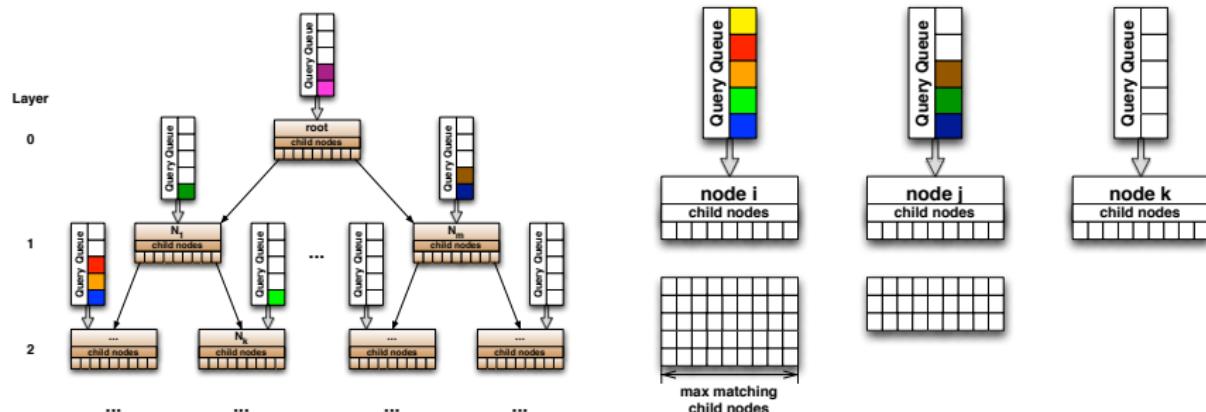
Scan Parallelization

Inter-Node Parallelization:

- Independent node batches
- Pipelining

Intra-Node Parallelization:

- Independent predicate tests
- SIMD features



GPU Implementation - Processing Model

- Nvidia CUDA implementation
- Multiple cores per GPU device
 - Execution of independent subtasks without synchronization
 - Subtask = node - query batch pair
- Multiple thread processors per GPU core
 - Execution of data parallel instructions by separate threads
 - Synchronization possible
 - Instruction = predicate test

GPU Implementation - Memory Hierarchy

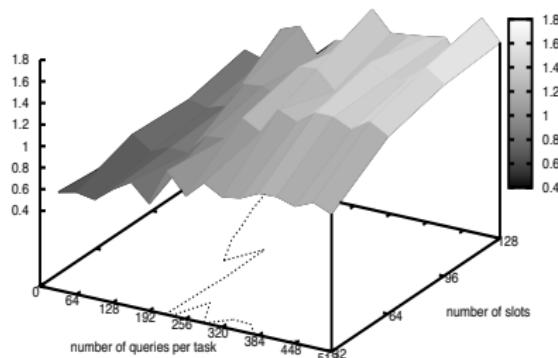
- Global main memory on device
 - Explicit transfer from host memory via PCIe bus
 - Caching of node and query data to avoid transfer overhead
 - Scan preparation phase to determine input data offsets
- Shared memory on each die
 - Two orders of magnitude faster than global memory
 - Software-controlled cache for scan data

Setup

- 3-D R-tree implementation
 - Generated index nodes
 - Intel Xeon X5690 CPU
 - Nvidia Tesla C2050 GPU
-
- Where shall a scan be executed?
 - How can the performance be improved with hybrid processing?

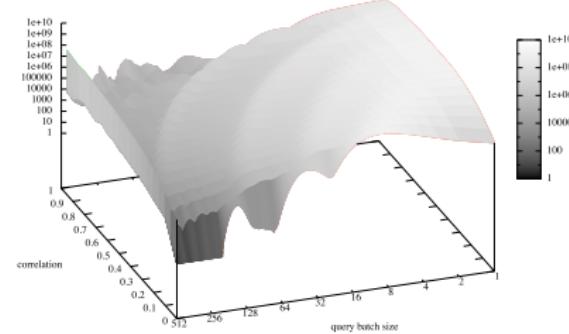
Tree Parameters

- Generated index nodes
- Generated predicates
- Full processor utilization
- Overheads included for GPU measurements
- $speedup = \frac{CPU\ time}{GPU\ time}$



Workload Simulation

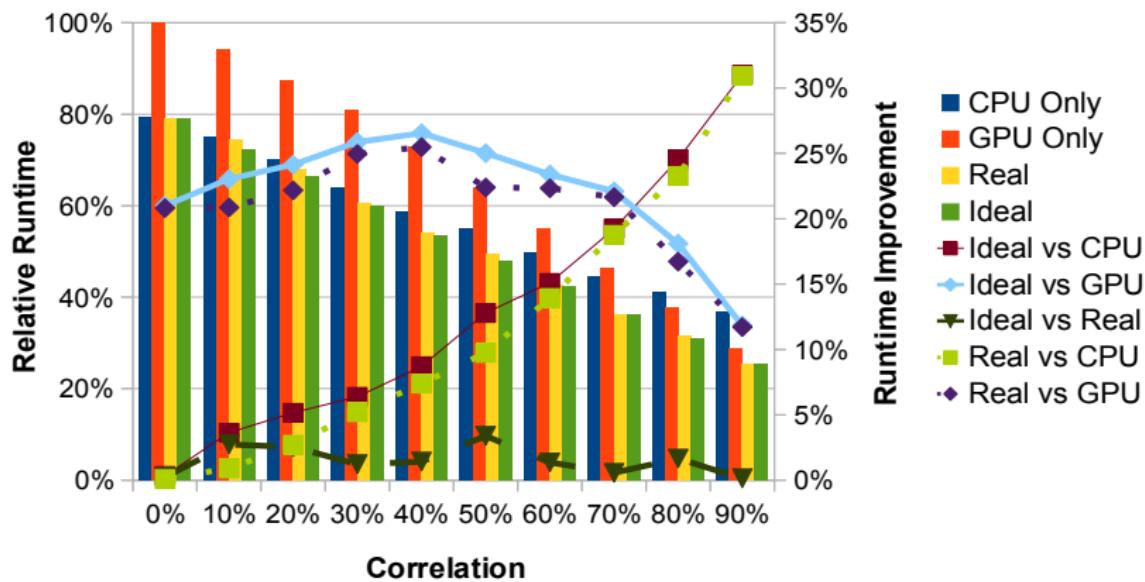
- CPU for small batches
- GPU for large batches
- How do batch sizes change when queries are streamed through the tree?
- Simulation for full R-tree
 - 96 slots per node
 - 5 layers → 8 billion indexed entries
 - 10.000 root queries



Workload Simulation - Parameter Correlation

Runtime Improvements Index Scan

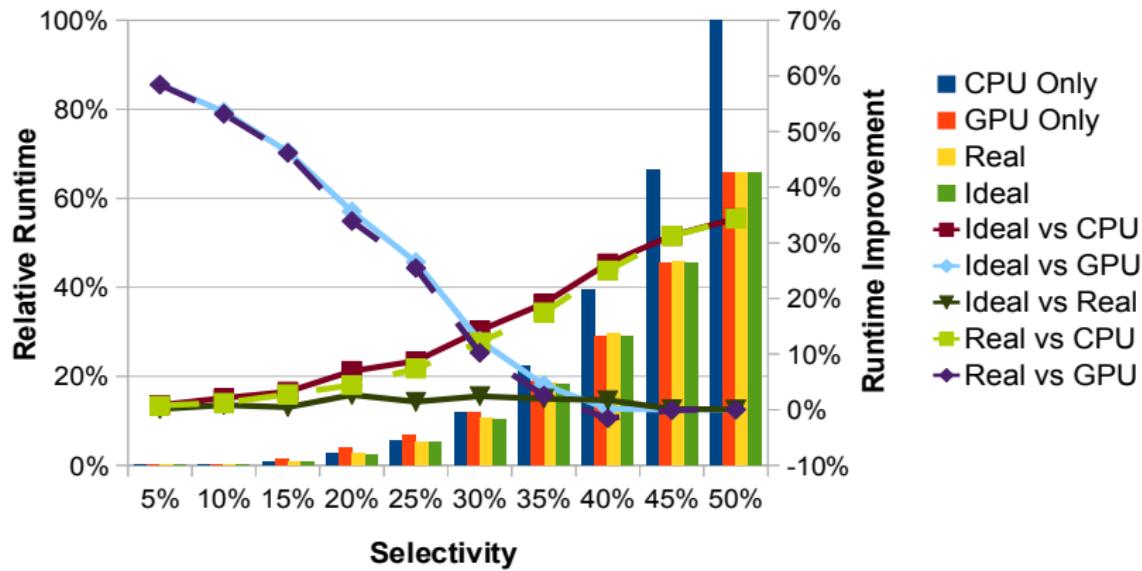
96 Slots, 25% Selectivity, 10000 Root Queries, 512 MaxBatchSize



Workload Simulation - Parameter Selectivity

Runtime Improvements Index Scan

96 Slots, 25% Duplicates, 10000 Root Queries, 512 MaxBatchSize



Conclusion & Outlook

Conclusion

- Extended GiST with hardware abstraction layer
- Performance improvements are possible
- Overheads are not negligible!

Next steps

- Prototype improvements
- Specialization for other tree types
- Full support for all GiST operations

References

- [1] http://en.wikipedia.org/wiki/File:Ray_trace_diagram.svg.
- [2] Science and Technology Review, June 2007.
"Virtual Dams Subjected to Strong Earthquakes".
- [3] CHOURASIA, A., OLSEN, K., CUI, Y., LEE, K., ZHOU, J., ELY, G., SMALL, P., ROTEN, D., DAY, S., MAECHLING, P., JORDAN, T., PANDA, D. K., AND LEVESQUE, J.
Ground motion visualization of M8 earthquake simulation using height field.
In *SciDAC* (2011).
Available at <http://www.mcs.anl.gov/uploads/cels/papers/scidac11/>.
- [4] HELLERSTEIN, J. M., NAUGHTON, J. F., AND PFEFFER, A.
Generalized Search Trees for Database Systems.
In *VLDB* (1995).
- [5] KAVAN, L., AND ZARA, J.
Fast Collision Detection for Skeletally Deformable Models.
Computer Graphics Forum 24, 3 (2005), 363–372.

Discussion

Thanks for your attention!

Questions?